

1

# Autoencoders

## Data Science in Electron Microscopy

---

Philipp Pelz


2024

[https://github.com/ECLIPSE-Lab/WS24\\_DataScienceForEM](https://github.com/ECLIPSE-Lab/WS24_DataScienceForEM)



# What is an Autoencoder?

- Neural network architecture that learns to:
  - Compress (encode) data into a lower-dimensional representation
  - Reconstruct (decode) the original data from this representation
- Trained to minimize reconstruction error
- Learns efficient data representations unsupervised



Basic Autoencoder Architecture

# Autoencoder Components

- **Encoder:** Compresses input into latent representation
- **Latent Space:** Compressed representation of the data
- **Decoder:** Reconstructs input from latent representation
- Training objective: minimize difference between input and output

```

1 import torch
2 import torch.nn as nn
3
4 class ConvAutoencoder(nn.Module):
5     def __init__(self):
6         super().__init__()
7
8         # Encoder
9         self.encoder = nn.Sequential(
10             nn.Conv2d(1, 16, 3, stride=2, padding=1), # [B, 1, 28, 28] -> [B, 16, 14, 14]
11             nn.ReLU(),
12             nn.Conv2d(16, 32, 3, stride=2, padding=1), # [B, 16, 14, 14] -> [B, 32, 7, 7]
13             nn.ReLU(),
14             nn.Conv2d(32, 64, 7) # [B, 32, 7, 7] -> [B, 64, 1, 1]
15         )
16
17         # Decoder
18         self.decoder = nn.Sequential(
19             nn.ConvTranspose2d(64, 32, 7) # [B, 64, 1, 1] -> [B, 32, 7, 7]

```



# Training an Autoencoder

```
1 def train_autoencoder(model, train_loader, num_epochs=10):  
2     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
3     model = model.to(device)  
4     criterion = nn.MSELoss()  
5     optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)  
6  
7     for epoch in range(num_epochs):  
8         for data in train_loader:  
9             img = data[0].to(device)  
10  
11             # Forward pass  
12             output = model(img)  
13             loss = criterion(output, img)  
14  
15             # Backward pass  
16             optimizer.zero_grad()  
17             loss.backward()  
18             optimizer.step()
```

# Applications of Autoencoders

- **Dimensionality Reduction**
  - Alternative to PCA
  - Can capture non-linear relationships
- **Denoising**
  - Train to reconstruct clean data from noisy input
  - Useful for image restoration
- **Feature Learning**
  - Learn meaningful representations for downstream tasks
  - Transfer learning

# Variations of Autoencoders

- **Denoising Autoencoders**
  - Add noise to input during training
  - Learn to recover original data
- **Variational Autoencoders (VAE)**
  - Learn probabilistic encodings
  - Generate new samples
- **Sparse Autoencoders**
  - Add sparsity constraints to latent representation
  - Learn more efficient encodings

# Example: Denoising Autoencoder

```
1 def add_noise(img, noise_factor=0.3):
2     noisy = img + noise_factor * torch.randn(*img.shape)
3     return torch.clamp(noisy, 0., 1.)
4
5 def train_denoising_autoencoder(model, train_loader, num_epochs=10):
6     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
7     model = model.to(device)
8     criterion = nn.MSELoss()
9     optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
10
11     for epoch in range(num_epochs):
12         for data in train_loader:
13             img = data[0].to(device)
14             noisy_img = add_noise(img)
15
16             # Forward pass
17             output = model(noisy_img)
18             loss = criterion(output, img) # Compare with clean image
19
```

# Practical Tips for Autoencoders

- Choose appropriate architecture for your data type
  - CNNs for images
  - RNNs for sequences
  - Dense layers for tabular data
- Consider:
  - Latent space dimension
  - Depth of encoder/decoder
  - Loss function
  - Regularization techniques
- Common issues:
  - Overfitting
  - Underfitting
  - Mode collapse (in VAEs)
  - Reconstruction quality vs. compression trade-off



# Variational Autoencoders (VAEs)

- Extension of traditional autoencoders that learns a **probabilistic** latent representation
- Instead of encoding to fixed points, encodes to probability distributions
- Enables:
  - Principled generation of new samples
  - Meaningful latent space interpolation
  - Better regularization of the latent space

# VAE vs. Traditional Autoencoder

## Traditional Autoencoder

- Deterministic encoding
- Point-wise latent representation
- No guarantee of continuous latent space
- Focus on reconstruction

## Variational Autoencoder

- Probabilistic encoding
- Distribution-based latent representation
- Continuous, structured latent space
- Balance between reconstruction and regularization

# VAE Mathematics

Instead of encoding input  $x$  to a point, VAE encodes to parameters of a distribution:

- Encoder outputs  $\mu$  and  $\log\sigma^2$  for each latent dimension
- Latent vector is sampled:  $z = \mu + \sigma \odot \epsilon$ , where  $\epsilon \sim \mathcal{N}(0, I)$

The VAE loss has two terms:

$$\mathcal{L}_{\text{VAE}} = \mathcal{L}_{\text{reconstruction}} + \beta \cdot \mathcal{L}_{\text{KL}}$$

where:

$$\mathcal{L}_{\text{KL}} = \frac{1}{2} \sum_{j=1}^J (\mu_j^2 + \sigma_j^2 - \log(\sigma_j^2) - 1)$$



# VAE Implementation

```
1 class ConvVAE(nn.Module):
2     def __init__(self, latent_dim=32):
3         super().__init__()
4
5         # Encoder
6         self.encoder = nn.Sequential(
7             nn.Conv2d(1, 32, 3, stride=2, padding=1), # 28x28 -> 14x14
8             nn.ReLU(),
9             nn.Conv2d(32, 64, 3, stride=2, padding=1), # 14x14 -> 7x7
10            nn.ReLU(),
11            nn.Flatten(),
12            nn.Linear(64 * 7 * 7, 256)
13        )
14
15        # Latent space
16        self.fc_mu = nn.Linear(256, latent_dim)
17        self.fc_var = nn.Linear(256, latent_dim)
18
19        # Decoder
```

# Training a VAE

```
1 def vae_loss(recon_x, x, mu, log_var, beta=1.0):  
2     # Reconstruction loss (binary cross entropy)  
3     BCE = F.binary_cross_entropy(recon_x, x, reduction='sum')  
4  
5     # KL divergence loss  
6     KLD = -0.5 * torch.sum(1 + log_var - mu.pow(2) - log_var.exp())  
7  
8     return BCE + beta * KLD
```



```
1 def train_vae(model, train_loader, num_epochs=10):
2     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3     model = model.to(device)
4     optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
5
6     for epoch in range(num_epochs):
7         for data in train_loader:
8             img = data[0].to(device)
9
10            # Forward pass
11            recon_batch, mu, log_var = model(img)
12            loss = vae_loss(recon_batch, img, mu, log_var)
13
14            # Backward pass
15            optimizer.zero_grad()
16            loss.backward()
17            optimizer.step()
```

# VAE Latent Space Properties

- **Continuous:** Similar points in latent space decode to similar images
- **Structured:** Enforced by KL divergence term
- **Meaningful:** Can perform interpolation and arithmetic in latent space

VAE Latent Space Visualization

# Generating New Samples with VAE

```
1 def generate_samples(model, num_samples=1):
2     with torch.no_grad():
3         # Sample from standard normal distribution
4         z = torch.randn(num_samples, model.latent_dim).to(device)
5         # Decode the samples
6         samples = model.decode(z)
7     return samples
8
9 def interpolate(model, img1, img2, steps=10):
10    # Encode both images
11    mu1, _ = model.encode(img1)
12    mu2, _ = model.encode(img2)
13
14    # Create interpolation points
15    alphas = torch.linspace(0, 1, steps)
16    interpolated = []
17
18    with torch.no_grad():
19        for alpha in alphas:
```



# Key Differences Summary

## 1. Latent Space

- Vanilla: Discrete, potentially discontinuous
- VAE: Continuous, probabilistic

## 2. Loss Function

- Vanilla: Only reconstruction loss
- VAE: Reconstruction + KL divergence loss

## 3. Generation Capabilities

- Vanilla: Limited/unreliable
- VAE: Principled generation of new samples

## 4. Training Stability

- Vanilla: Can be unstable
- VAE: More stable due to regularization

# Example training a VAE

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from torch.utils.data import DataLoader
5 from torchvision import datasets, transforms
6 import matplotlib.pyplot as plt
7
8
9 def add_noise(x, noise_factor=0.3):
10     noisy = x + noise_factor * torch.randn_like(x)
11     return torch.clamp(noisy, 0., 1.)
12
13 def train_epoch(model, dataloader, optimizer, device, noise_factor=0.3):
14     model.train()
15     train_loss = 0
16
17     for batch_idx, (data, _) in enumerate(dataloader):
18         data = data.to(device)
19         noisy_data = add_noise(data, noise_factor)
```

Using device: cuda

Training completed and model saved!

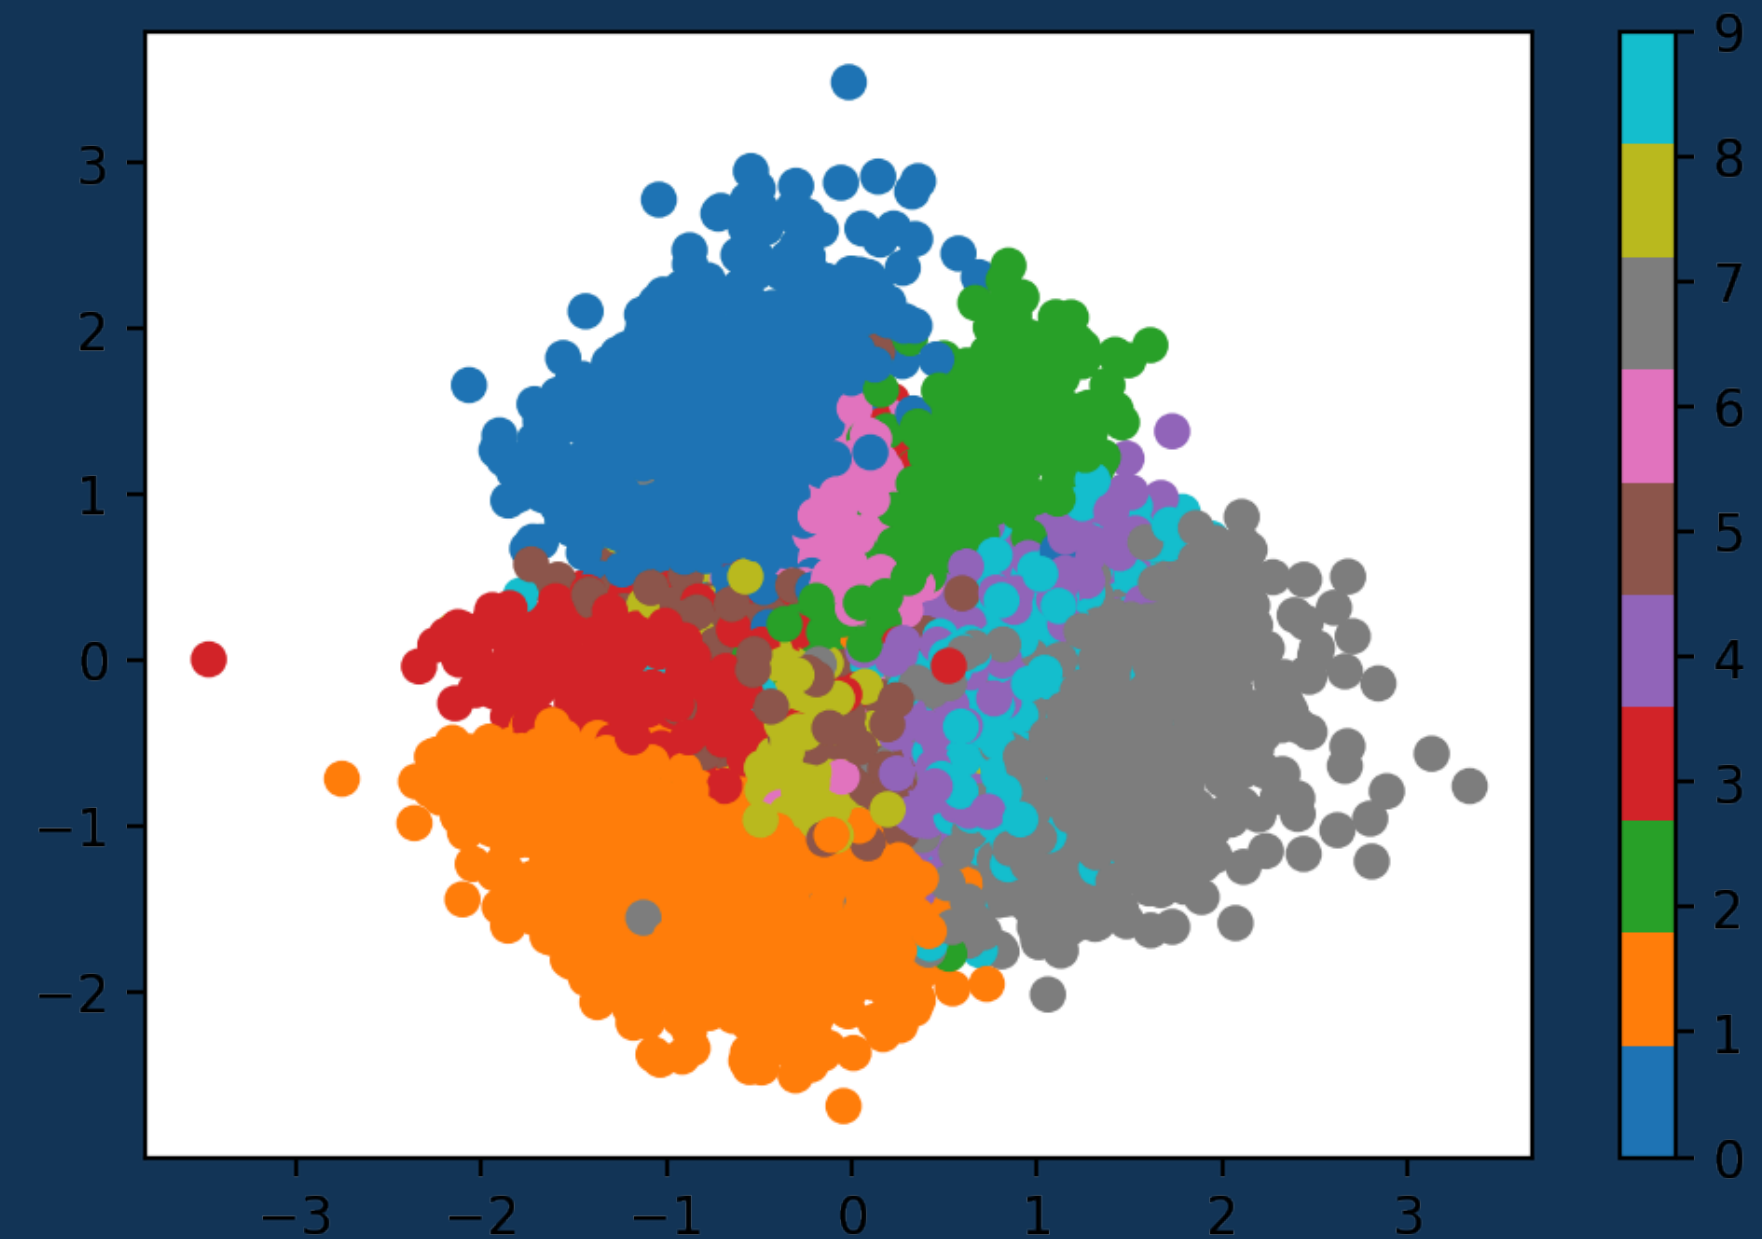
# VAE vs. Traditional Autoencoder

## Traditional Autoencoder



Vanilla AE Latent Space Visualization

## Variational Autoencoder



VAE Latent Space Visualization